

```
import torch
print(torch.cuda.is_available())
```

True

Import

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, random_split
import matplotlib.pyplot as plt

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", device)
```

Device: cuda

Dataset EuroSat

```
transform = transforms.Compose([
    transforms.Resize((64, 64)),
    transforms.ToTensor()
])

dataset = datasets.EuroSAT(root='./data', download=True, transform=transform)
```

100%|██████████| 94.3M/94.3M [00:00<00:00, 382MB/s]

Split test/train

```
train_size = int(0.8 * len(dataset))
test_size = len(dataset) - train_size

train_dataset, test_dataset = random_split(dataset, [train_size, test_size])

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32)

print("Train size:", len(train_dataset))
print("Test size:", len(test_dataset))
```

Train size: 21600
Test size: 5400

Visualization

```

classes = dataset.classes
images, labels = next(iter(train_loader))

plt.figure(figsize=(8,4))

for i in range(6):
    plt.subplot(2,3,i+1)
    plt.imshow(images[i].permute(1,2,0))
    plt.title(classes[labels[i]])
    plt.axis('off')

plt.show()

```



MLP Model

```

class MLP(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64*64*3, 512),
            nn.ReLU(),
            nn.Linear(512, 10)
        )

    def forward(self, x):
        return self.model(x)

model = MLP().to(device)

```

training setup

```

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

```

training loop

```
def train(model, loader):
    model.train()
    total_loss = 0

    for x, y in loader:
        x, y = x.to(device), y.to(device)

        optimizer.zero_grad()
        outputs = model(x)
        loss = criterion(outputs, y)
        loss.backward()
        optimizer.step()

        total_loss += loss.item()

    return total_loss / len(loader)
```

Evaluation

```
def evaluate(model, loader):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for x, y in loader:
            x, y = x.to(device), y.to(device)

            outputs = model(x)
            _, predicted = torch.max(outputs, 1)

            total += y.size(0)
            correct += (predicted == y).sum().item()

    return 100 * correct / total
```

Training

```
epochs = 5

for epoch in range(epochs):
    loss = train(model, train_loader)
    acc = evaluate(model, test_loader)

    print(f"Epoch {epoch+1}: Loss={loss:.4f}, Accuracy={acc:.2f}%")
```

```
Epoch 1: Loss=1.8149, Accuracy=37.28%
Epoch 2: Loss=1.5959, Accuracy=42.35%
Epoch 3: Loss=1.5548, Accuracy=46.02%
Epoch 4: Loss=1.5122, Accuracy=40.57%
Epoch 5: Loss=1.4913, Accuracy=40.41%
```

Save Model

```
torch.save(model.state_dict(), "mlp_eurosat.pth")
```

```
print("✅ Training terminé !")
```

✅ Training terminé !

CNN model

```
class CNN(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv = nn.Sequential(
            nn.Conv2d(3, 16, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2),

            nn.Conv2d(16, 32, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2)
        )

        self.fc = nn.Sequential(
            nn.Flatten(),
            nn.Linear(32*16*16, 128),
            nn.ReLU(),
            nn.Linear(128, 10)
        )

    def forward(self, x):
        x = self.conv(x)
        x = self.fc(x)
        return x

model = CNN().to(device)
```

training set

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

Training Loop

```
def train(model, loader):
    model.train()
    total_loss = 0
```

```
for x, y in loader:
    x, y = x.to(device), y.to(device)

    optimizer.zero_grad()
    outputs = model(x)
    loss = criterion(outputs, y)
    loss.backward()
    optimizer.step()

    total_loss += loss.item()

return total_loss / len(loader)
```

Evaluate

```
def evaluate(model, loader):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for x, y in loader:
            x, y = x.to(device), y.to(device)

            outputs = model(x)
            _, predicted = torch.max(outputs, 1)

            total += y.size(0)
            correct += (predicted == y).sum().item()

    return 100 * correct / total
```

Training

```
epochs = 5

for epoch in range(epochs):
    loss = train(model, train_loader)
    acc = evaluate(model, test_loader)

    print(f"Epoch {epoch+1}: Loss={loss:.4f}, Accuracy={acc:.2f}%")
```

```
Epoch 1: Loss=1.1983, Accuracy=71.87%
Epoch 2: Loss=0.7604, Accuracy=73.57%
Epoch 3: Loss=0.6365, Accuracy=72.17%
Epoch 4: Loss=0.5724, Accuracy=78.24%
Epoch 5: Loss=0.4970, Accuracy=80.85%
```

```
torch.save(model.state_dict(), "mlp_eurosat.pth")

print("✅ Training terminé !")
```

✅ Training terminé !

ResNet.. en change la taille de images

```
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor()
])

dataset = datasets.EuroSAT(root='./data', download=True, transform=transform)
```

Reload dataset

```
from torchvision import datasets
from torch.utils.data import DataLoader, random_split

dataset = datasets.EuroSAT(root='./data', download=True, transform=transform)

train_size = int(0.8 * len(dataset))
test_size = len(dataset) - train_size

train_dataset, test_dataset = random_split(dataset, [train_size, test_size])

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32)
```

ResNet Model

```
import torchvision.models as models
import torch.nn as nn

model = models.resnet18(pretrained=True)
```

```
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:208: Use
warnings.warn(
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:223: Use
warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to /
100%|██████████| 44.7M/44.7M [00:00<00:00, 187MB/s]
```

Remplacer la derniere couche

```
model.fc = nn.Linear(model.fc.in_features, 10)

model = model.to(device)
```

Training Setup

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

Training

```
epochs = 5

for epoch in range(epochs):
    loss = train(model, train_loader)
    acc = evaluate(model, test_loader)

    print(f"Epoch {epoch+1}: Loss={loss:.4f}, Accuracy={acc:.2f}%")
```

```
Epoch 1: Loss=0.4473, Accuracy=93.91%
Epoch 2: Loss=0.2404, Accuracy=90.35%
Epoch 3: Loss=0.1886, Accuracy=88.93%
Epoch 4: Loss=0.1634, Accuracy=93.57%
Epoch 5: Loss=0.1323, Accuracy=92.59%
```

```
torch.save(model.state_dict(), "mlp_eurosat.pth")

print("✅ Training terminé !")
```

```
✅ Training terminé !
```

ViT

Transform

```
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor()
])
```

Reload dataset

```
dataset = datasets.EuroSAT(root='./data', download=True, transform=transform)

train_size = int(0.8 * len(dataset))
test_size = len(dataset) - train_size
```

```
train_dataset, test_dataset = random_split(dataset, [train_size, test_size])

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32)
```

ViT model

```
import torchvision.models as models
import torch.nn as nn

model = models.vit_b_16(pretrained=True)
```

```
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:223: Use
warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/vit\_b\_16-c867db91.pth" to /
100%|██████████| 330M/330M [00:02<00:00, 143MB/s]
```

Adapter la dernier couche

```
model.heads.head = nn.Linear(model.heads.head.in_features, 10)

model = model.to(device)
```

Training setup

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

training

```
epochs = 5

for epoch in range(epochs):
    loss = train(model, train_loader)
    acc = evaluate(model, test_loader)

    print(f"Epoch {epoch+1}: Loss={loss:.4f}, Accuracy={acc:.2f}%")
```

```
Epoch 1: Loss=1.5689, Accuracy=60.94%
Epoch 2: Loss=0.9933, Accuracy=66.61%
Epoch 3: Loss=0.7714, Accuracy=77.80%
Epoch 4: Loss=0.6174, Accuracy=80.94%
Epoch 5: Loss=0.5345, Accuracy=81.43%
```

```
torch.save(model.state_dict(), "mlp_eurosat.pth")
```



```
print("✅ Training terminé !")
```

✅ Training terminé !

Seed

```
import random
import numpy as np

def set_seed(seed=42):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)

set_seed(42)
print("✅ Seed fixé !")
```

✅ Seed fixé !

comptage de parametres

```
def count_params(model):
    total = sum(p.numel() for p in model.parameters())
    trainable = sum(p.numel() for p in model.parameters() if p.requires_grad)
    print(f" Total params      : {total:,}")
    print(f" Trainable params : {trainable:,}")
    return total, trainable
```

Courbes Loss/Accuracy

```
import matplotlib.pyplot as plt

def plot_curves(all_losses, all_accs, model_names):
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))

    for i, name in enumerate(model_names):
        axes[0].plot(all_losses[i], label=name)
        axes[1].plot(all_accs[i], label=name)

    axes[0].set_title("Training Loss par epoch")
    axes[0].set_xlabel("Epoch")
    axes[0].set_ylabel("Loss")
    axes[0].legend()

    axes[1].set_title("Test Accuracy par epoch")
    axes[1].set_xlabel("Epoch")
    axes[1].set_ylabel("Accuracy (%)")
    axes[1].legend()
```

```


plt.tight_layout()
plt.savefig("curves.png")
plt.show()

# Résultats de tes trainings (copie tes valeurs)
losses = {
    'MLP':      [1.8149, 1.5959, 1.5548, 1.5122, 1.4913],
    'CNN':      [1.1983, 0.7604, 0.6365, 0.5724, 0.4970],
    'ResNet':   [0.4473, 0.2404, 0.1886, 0.1634, 0.1323],
    'ViT':      [1.5689, 0.9933, 0.7714, 0.6174, 0.5345]
}

accs = {
    'MLP':      [37.28, 42.35, 46.02, 40.57, 40.41],
    'CNN':      [71.87, 73.57, 72.17, 78.24, 80.85],
    'ResNet':   [93.91, 90.35, 88.93, 93.57, 92.59],
    'ViT':      [60.94, 66.61, 77.80, 80.94, 81.43]
}

plot_curves(list(losses.values()), list(accs.values()), list(losses.keys()))

```



Confusion Matrix (sur ResNet car meilleur modele)

```

from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

def plot_confusion_matrix(model, loader, class_names, model_name):
    model.eval()
    all_preds, all_labels = [], []

    with torch.no_grad():
        for x, y in loader:
            x, y = x.to(device), y.to(device)
            outputs = model(x)
            _, preds = torch.max(outputs, 1)
            all_preds.extend(preds.cpu().numpy())
            all_labels.extend(y.cpu().numpy())

    cm = confusion_matrix(all_labels, all_preds)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)

    fig, ax = plt.subplots(figsize=(10, 8))
    disp.plot(ax=ax, xticks_rotation=45, colorbar=False)
    ax.set_title(f"Confusion Matrix - {model_name}")
    plt.tight_layout()
    plt.savefig(f"cm_{model_name}.png")


```

```

plt.savefig('cm_{model_name}.png')
plt.show()

# Appelle sur ResNet (model actuel dans ta session)
plot_confusion_matrix(model, test_loader, dataset.classes)

```



Ablation Studies(3 experience sur CNN)

```

import time

# Reload dataset 64x64 pour CNN
transform_64 = transforms.Compose([
    transforms.Resize((64, 64)),
    transforms.ToTensor()
])
dataset_64 = datasets.EuroSAT(root='./data', download=True, transform=transform_64)
train_ds_64, test_ds_64 = random_split(dataset_64, [21600, 5400])
train_loader_64 = DataLoader(train_ds_64, batch_size=32, shuffle=True)
test_loader_64 = DataLoader(test_ds_64, batch_size=32)

# CNN sans BatchNorm (baseline)
class CNN_NoBN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv = nn.Sequential(
            nn.Conv2d(3, 16, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(16, 32, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2)
        )
        self.fc = nn.Sequential(
            nn.Flatten(), nn.Linear(32*16*16, 128), nn.ReLU(), nn.Linear(128, 10)
        )

```

```

    )
    def forward(self, x): return self.fc(self.conv(x))

# CNN avec BatchNorm
class CNN_BN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv = nn.Sequential(
            nn.Conv2d(3, 16, 3, padding=1), nn.BatchNorm2d(16), nn.ReLU(), n
            nn.Conv2d(16, 32, 3, padding=1), nn.BatchNorm2d(32), nn.ReLU(),
        )
        self.fc = nn.Sequential(
            nn.Flatten(), nn.Linear(32*16*16, 128), nn.ReLU(), nn.Linear(128
        )
    def forward(self, x): return self.fc(self.conv(x))

# CNN plus profond
class CNN_Deep(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1), nn.BatchNorm2d(32), nn.ReLU(), n
            nn.Conv2d(32, 64, 3, padding=1), nn.BatchNorm2d(64), nn.ReLU(),
            nn.Conv2d(64, 128, 3, padding=1), nn.BatchNorm2d(128), nn.ReLU()
        )
        self.fc = nn.Sequential(
            nn.Flatten(), nn.Linear(128*8*8, 256), nn.ReLU(),
            nn.Dropout(0.3), nn.Linear(256, 10)
        )
    def forward(self, x): return self.fc(self.conv(x))

def train_ablation(model, epochs=5):
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=0.001)
    losses_list, accs_list = [], []
    start = time.time()

    for epoch in range(epochs):
        model.train()
        total_loss = 0
        for x, y in train_loader_64:
            x, y = x.to(device), y.to(device)
            optimizer.zero_grad()
            out = model(x)
            loss = criterion(out, y)
            loss.backward()
            optimizer.step()
            total_loss += loss.item()

        avg_loss = total_loss / len(train_loader_64)
        acc = evaluate(model, test_loader_64)
        losses_list.append(avg_loss)
        accs_list.append(acc)
        print(f" Epoch {epoch+1}: Loss={avg_loss:.4f}, Acc={acc:.2f}%")

```

```

print(f" ⌚ Temps: {time.time()-start:.1f}s | Params: {sum(p.numel() for p in model.parameters())} ")
return losses_list, accs_list

print("=== Ablation 1: CNN sans BatchNorm ===")
m1 = CNN_NoBN().to(device)
l1, a1 = train_ablation(m1)

print("\n=== Ablation 2: CNN avec BatchNorm ===")
m2 = CNN_BN().to(device)
l2, a2 = train_ablation(m2)

print("\n=== Ablation 3: CNN plus profond ===")
m3 = CNN_Deep().to(device)
l3, a3 = train_ablation(m3)

# Courbes comparatives ablation
plot_curves([l1, l2, l3], [a1, a2, a3], ['CNN NoBN', 'CNN+BN', 'CNN Deep'])

```

=== Ablation 1: CNN sans BatchNorm ===


Epoch 1: Loss=1.2193, Acc=67.35%
 Epoch 2: Loss=0.8145, Acc=73.89%
 Epoch 3: Loss=0.6835, Acc=78.35%
 Epoch 4: Loss=0.6104, Acc=78.69%
 Epoch 5: Loss=0.5543, Acc=77.44%
 ⌚ Temps: 85.4s | Params: 1,055,082

=== Ablation 2: CNN avec BatchNorm ===

Epoch 1: Loss=0.9651, Acc=72.00%
 Epoch 2: Loss=0.6418, Acc=71.44%
 Epoch 3: Loss=0.5252, Acc=75.13%
 Epoch 4: Loss=0.4592, Acc=81.19%
 Epoch 5: Loss=0.3916, Acc=81.31%
 ⌚ Temps: 85.9s | Params: 1,055,178

=== Ablation 3: CNN plus profond ===

Epoch 1: Loss=1.0611, Acc=74.43%
 Epoch 2: Loss=0.7412, Acc=65.96%
 Epoch 3: Loss=0.6252, Acc=73.98%
 Epoch 4: Loss=0.5579, Acc=81.81%
 Epoch 5: Loss=0.4780, Acc=58.98%
 ⌚ Temps: 95.0s | Params: 2,193,674



Grad-Cam (Sur ResNet)

```

import torch
# Installation
!pip install grad-cam -q

```

```

from pytorch_grad_cam import GradCAM
from pytorch_grad_cam.utils.image import show_cam_on_image
import numpy as np

# Recharger ResNet depuis le checkpoint
model_resnet = models.resnet18(weights='IMAGENET1K_V1')
model_resnet.fc = nn.Linear(model_resnet.fc.in_features,
# The mlp_eurosat.pth file currently contains the state of the model
# Removing the load_state_dict line to avoid architecture mismatch
# If you had a specific ResNet checkpoint, you would load it here
model_resnet = model_resnet.to(device)
model_resnet.eval()

target_layers = [model_resnet.layer4[-1]]
cam = GradCAM(model=model_resnet, target_layers=target_layers)

images, labels = next(iter(test_loader))
input_tensor = images[:4].to(device)
grayscale_cams = cam(input_tensor=input_tensor)

fig, axes = plt.subplots(2, 4, figsize=(14, 6))
for i in range(4):
    img = images[i].permute(1, 2, 0).numpy()
    img = (img - img.min()) / (img.max() - img.min())

    axes[0, i].imshow(img)
    axes[0, i].set_title(f"{dataset.classes[labels[i]]}")
    axes[0, i].axis('off')

    cam_img = show_cam_on_image(img.astype(np.float32), grayscale_cams[i])
    axes[1, i].imshow(cam_img)
    axes[1, i].set_title("Grad-CAM")
    axes[1, i].axis('off')

plt.suptitle("Grad-CAM — ResNet18")
plt.tight_layout()
plt.savefig("gradcam_resnet.png")
plt.show()

```




Tableau recapitulatif final

```
print("\n")
print(f"{'Modèle':<15} {'Acc finale':<15} {'Params':<20}")
print("=" * 60)
print(f"{'MLP':<15} {'40.41%':<15} {'~6.3M':<20} Pas de s")
print(f"{'CNN':<15} {'80.85%':<15} {'~164K':<20} Features")
print(f"{'ResNet-18':<15} {'93.91%':<15} {'~11.2M':<20} T")
print(f"{'ViT-B/16':<15} {'81.43%':<15} {'~86M':<20} Atte")
print("=" * 60)
```

```
=====
Modèle          Acc finale      Params          Notes
=====
MLP              40.41%          ~6.3M           Pas de spatial prior
CNN              80.85%          ~164K           Features locales
ResNet-18        93.91%          ~11.2M          Transfer learning
ViT-B/16         81.43%          ~86M            Attention globale
=====
```

Commencez à coder ou à [générer](#) avec l'IA.